

Continual Learning with O-LoRA: Orthogonal Low-Rank Adapters for Sequential Tasks

Alex Birsan, Radu Savin

Coordinators: Marius Ninel Cătălin Drăgoi, Florin Brad
University of Bucharest, Romania



UNIVERSITY OF
BUCHAREST
— VIRTUTE ET SAPIENTIA

1. Introduction

Continual learning requires a model to sequentially acquire knowledge across multiple tasks without forgetting what it previously learned. A common challenge is **catastrophic forgetting**: adapting to a new task overwrites representations built for earlier ones.

When parameter-efficient methods such as **LoRA** are applied in this setting, the low-rank adapter matrices trained on successive tasks tend to occupy overlapping subspaces, causing destructive interference.

O-LoRA addresses this by imposing an orthogonality constraint on the A matrices of each new task adapter with respect to all previously learned ones, ensuring each task projects inputs into a distinct region of the low-rank latent space and limiting the interference that drives forgetting.

2. Dataset and Preprocessing

Task sequence (fixed order):

- **AG News** – 4-class news topic classification
- **Amazon Polarity** – binary sentiment classification
- **DBpedia 14** – 14-class Wikipedia topic classification
- **Yahoo Answers Topics** – 10-class question topics

Each task is trained on independently, with no replay of prior task data.

Instruction-tuning prompt format:

```
{task instruction}\n
Option: {comma-separated label names}\n
{input text}\n
Answer:
```

Only the input text is truncated when the sequence exceeds the maximum length. Cross-entropy loss is computed only on the label token; at evaluation the predicted label is the candidate label token with maximum logit at the position before **Answer**.

3. Models

Backbone: Qwen2.5-1.5B. **LoRA:** rank $r=8$, $\alpha=32$, dropout 0.1, on Q, V of every attention layer. **Optim:** 1 epoch/task, lr 10^{-3} , batch 8.

Compared methods:

- **IncLoRA** – separate frozen adapter per task (BWT = 0 by construction).
- **O-LoRA** – one cumulative adapter; orthogonality penalty between the new task’s A and all previous ones, weighted by $\lambda_1 \in [0.2, 0.3]$.

Orthogonal loss. For task t with $\Delta W_t = B_t A_t$:

$$\mathcal{L}_t = \mathcal{L}_{\text{CE}}(A_t, B_t) + \lambda_1 \sum_{i < t} \|A_t A_i^\top\|_F^2$$

The penalty drives row-spaces of successive A into orthogonal directions. Crucially, **no constraint on B** .

4. Results

IncLoRA baseline (no forgetting):

Task	Accuracy
AG News	0.854
Amazon Polarity	0.956
DBpedia 14	0.972
Yahoo Answers Topics	0.640

O-LoRA – summary metrics:

Run	Samples / task	λ_1	Avg Acc	BWT
Short	500/500/500/500	0.3	0.846	−0.032
2h	500/650/1850/1300	0.2	0.799	−0.091
Full (~50%)	2100/2700/7500/5300	0.3	0.654	−0.268


Per-task accuracy: initial → final

Task	Short	2h	Full (~50%)
AG News	0.878→0.825 (−6%)	0.842→0.747 (−11%)	0.887→0.494 (−44%)
Amazon Polarity	0.956→0.957 (~0%)	0.947→0.915 (−3%)	0.905→0.746 (−18%)
DBpedia 14	0.964→0.919 (−5%)	0.979→0.833 (−15%)	0.980→0.729 (−26%)
Yahoo Answers Topics	0.683→0.683 (—)	0.702→0.702 (—)	0.649→0.649 (—)

Avg accuracy and BWT both **degrade as the number of training samples grows**, even though more data improves *initial* per-task accuracy. The **instruction + options** prompt format yields +0.20 Avg Acc and +0.25 BWT – a larger effect than any hyperparameter choice.

5. Orthogonality Analysis

Three pairwise metrics are computed across all task adapter matrices: cosine similarity (flat-vector alignment), Gram Frobenius norm $/r$ (the quantity directly minimised by O-LoRA’s loss term), and max principal angle cosine (closest pair of subspace directions).



Task-pair similarity of A matrices (8h): cosine, Gram Frobenius norm $/r$, max principal angle cosine – **all near-zero off-diagonal**: A matrices are geometrically separated.

Task-pair similarity of the full products $\Delta W = BA$ (8h) – **substantially higher off-diagonal values**: orthogonality on A does not transfer to BA .

All three metrics show near-zero similarity across A even with $\lambda_1 \in [0.2, 0.3]$ (lower than in the original paper) – the constraint is easy to satisfy. However B matrices, unconstrained by the loss, share $\sim 30\%$ of their column-space directions (principal-angle cosine ~ 0.29) and grow in magnitude with training. As a consequence, the *combined* BA products lose orthogonality (right figure) – **plausibly the primary driver of forgetting at higher scales**.

6. Conclusion

- O-LoRA’s regulariser keeps A matrices **geometrically separated** across all three metrics with mild penalties ($\lambda_1 \in [0.2, 0.3]$, lower than in the original paper) – the constraint is easy to satisfy.
- Translating geometric separation into reduced forgetting is **highly sensitive to data scale**: near-zero BWT (−0.032) at 500 samples/task, sharp collapse (−0.268) at $\sim 50\%$ of the paper’s training scale.
- The unregularised B matrices grow in magnitude and share $\sim 30\%$ of their column-space directions across tasks, so the combined products BA lose orthogonality – a likely primary driver of forgetting at scale.
- **Prompt structure matters more than hyperparameters**: listing label names in every prompt yields +0.20 Avg Acc and +0.25 BWT, the largest single improvement we observed.
- Non-monotonic recovery of AG News across the 8h run is consistent with unconstrained B matrices coincidentally cancelling each other’s interference at intermediate steps, rather than structured preservation of past knowledge.
- **Future work**: regularise B as well, or constrain $\|BA\|$ directly rather than $\|A_t A_i^\top\|$.